

Historical Record: career-transition-intake and career-transition (Pre-Merger)

Prepared: 2026-08-01 Status: Complete record of both repos' independent history, before they were merged into a single monorepo (career-transition, career-transition-intake#38). Covers everything that happened while each was its own separate GitHub repository — commits, every issue and every comment on it, pull requests, and releases. Work that happened after the merge (the migration follow-through, the intake form restyle, the copy repositioning) is tracked in the new repo's own issues (#8, #9) and is out of scope here; this document is specifically the record of how each project got to the point of being merged.

Prepared so the original career-transition-intake GitHub repo (currently archived, read-only) and the local pre-merge backup directories can eventually be deleted without losing anything — once this document is confirmed accurate.

Part 1: career-transition-intake — the Flask intake app

Formation and evolution

The intake app began life as a Gradio application on Hugging Face Spaces — the fastest way to get a working onboarding form in front of the first client. As real usage started, that platform choice became a bottleneck, and the project moved through Google Cloud Run before settling on Render, with Gradio itself eventually replaced entirely by a hand-built Flask application with a Bootstrap 5 (dark-themed, at the time) form — the shape the app kept for most of its life. Email delivery was switched from SMTP (blocked on Render's outbound port 587) to the Resend HTTP API early on, a decision that held for the rest of the project.

From there the app grew through several distinct phases visible in its own issue history: a security-hardening phase (two full audits, S-01 through S-19, covering CSRF, upload validation, rate limiting, log-injection, formula-injection in the exported metrics spreadsheet, and more); a UX-refinement phase driven by rounds of real user feedback (client-type conditional logic, a CV-or-background-questions fallback requirement added after two real intakes came in with neither, dynamic portfolio links replacing a free-text box); an operations phase (semantic versioning and a CHANGELOG.md, submission metrics extracted from Render's logs into a tracked spreadsheet, a branded cold-start loading page built to replace Render's own server-log screen during the free tier's slow wake-up); and finally a migration phase — a spike into collecting payment at submission time (M-Pesa/Pesapal), a spike into unifying the app with its companion landing page under one domain, a deep dive that turned into an actual migration off Render onto Vercel (Fluid Compute, Upstash Redis-backed rate limiting, a real production cutover), and the eventual decommissioning of Render entirely — the last major body of work before the two repos were merged into one.

Below: the complete commit history, every issue with its full comment thread, and every release, in full.

Complete commit history

85 commits, oldest to newest:

- fa58da12 (2026-07-09, ericgitonga) — Initial commit — onboarding form, HF deploy workflow
- 36a56c23 (2026-07-09, ericgitonga) — Move client data under Clients/, update .gitignore and SKILL.md

- bca30b68 (2026-07-09, ericgitonga) — Remove hosting.md from repo, add to .gitignore
- d47a6250 (2026-07-09, ericgitonga) — add generate_hosting_pdf.py to .gitignore
- 6310f7a1 (2026-07-09, ericgitonga) — Fix invalid HF Space colorTo value (teal → indigo)
- 7e4c88c3 (2026-07-09, ericgitonga) — Pin huggingface_hub<0.25.0 to fix gradio HfFolder import error
- 539ba22f (2026-07-09, ericgitonga) — Upgrade to gradio 5.50.0 and fix css/theme placement
- 852f9e42 (2026-07-09, ericgitonga) — Switch hosting from Hugging Face Spaces to Google Cloud Run
- 963c6517 (2026-07-09, ericgitonga) — Switch hosting from Google Cloud Run to Render
- 08fabcb6f (2026-07-09, ericgitonga) — Downgrade to gradio 4.44.1 to fix blank screen on Render
- a5aa7171 (2026-07-09, ericgitonga) — Pin Python to 3.11.9 to fix pydub/audioop crash on Render
- 6eb44736 (2026-07-09, ericgitonga) — Fix crash on Render: disable API info and pin gradio-client<1.0
- ffc6cf9d (2026-07-09, ericgitonga) — Remove conflicting gradio-client pin (gradio 4.44.1 requires ==1.3.0)
- 8f9722b7 (2026-07-09, ericgitonga) — Patch gradio_client boolean schema bug that crashes every request
- 29d314a7 (2026-07-09, ericgitonga) — Pin fastapi==0.115.0 to get compatible starlette (0.38.x)
- 5e4fb060 (2026-07-09, ericgitonga) — Bypass Gradio localhost check that always fails on Render
- de53d0a4 (2026-07-09, ericgitonga) — Replace Gradio with Flask
- c4a29af0 (2026-07-09, ericgitonga) — Fix SMTP worker timeout; add App Password setup guide
- 9e007621 (2026-07-09, ericgitonga) — Switch email sending from SMTP to Resend API
- 9ab0104e (2026-07-09, ericgitonga) — Remove Email Settings section; apply dark theme
- 82383df2 (2026-07-09, ericgitonga) — Update header text; add multi-file upload field
- 5306675b (2026-07-09, ericgitonga) — Remove headshot upload field
- 7b18c00c (2026-07-10, ericgitonga) — Add hosting.pdf to repo; update docs for Flask + Resend stack
- 2d7a1328 (2026-07-10, ericgitonga) — Add detailed docstrings to all functions in app.py and generate_hosting_pdf.py
- fdd986a6 (2026-07-10, ericgitonga) — Add generate_design_pdf.py — design process report generator
- 7773cd51 (2026-07-10, ericgitonga) — Remove Gradio traces; update SKILL.md for hosted Flask form
- caf552d6 (2026-07-10, ericgitonga) — Add generate_selfhosting_pdf.py — self-hosting guide generator
- b9bc80e5 (2026-07-10, ericgitonga) — Add detailed docstrings to all functions in generate_design_pdf.py
- 980f0261 (2026-07-10, ericgitonga) — Ignore generate_security_pdf.py — private audit script
- 5d4301d9 (2026-07-10, ericgitonga) — Security hardening: all 15 audit findings resolved
- 8485fc47 (2026-07-10, ericgitonga) — Auto-generate SECRET_KEY locally when env var is absent
- c8a2ae6f (2026-07-10, ericgitonga) — Update design process report with security content
- 22517ab8 (2026-07-10, ericgitonga) — add extras folder
- 2ba1e8c2 (2026-07-10, ericgitonga) — Apply Round 1 feedback: form flow, new question, UX copy, error handling
- 4d207314 (2026-07-11, ericgitonga) — Add clean loading screen to replace Render cold-start page
- 5736c2f8 (2026-07-11, ericgitonga) — Fix static loading site: move to loading/ subdirectory
- d6a11916 (2026-07-11, ericgitonga) — edit user facing message
- ffc89d1e (2026-07-11, ericgitonga) — Document user-facing behaviour rules in SKILL.md

- 12650ade (2026-07-11, ericgitonga) — Update design_process.pdf with loading screen UX development
- 8bf4c0a0 (2026-07-11, ericgitonga) — Add cold-start loading screen feedback to feedback_01.pdf
- 19349523 (2026-07-11, ericgitonga) — Untrack private PDF generators and report files
- 61bc1a23 (2026-07-11, ericgitonga) — Name intake files by client slug and bundle uploads into a ZIP
- 693c5956 (2026-07-12, ericgitonga) — specify expertise
- 7c4ec239 (2026-07-14, ericgitonga) — Expand intake form with client segmentation, CV fallback, and work-style fields
- 1cb141ad (2026-07-14, ericgitonga) — Fix CV fallback position and add free-form notes field (closes #13)
- 1b6bf913 (2026-07-14, ericgitonga) — Replace portfolio textarea with dynamic link fields (closes #14)
- fc9b16e9 (2026-07-14, ericgitonga) — Instruct plan generator to fetch portfolio links via WebFetch (closes #15)
- bc6b1d5b (2026-07-14, ericgitonga) — Add client-type conditional logic and improve org stage descriptions (closes #16)
- c9bb4f68 (2026-07-14, ericgitonga) — Show clarification field when Other is selected in client type (closes #17)
- 3d79d419 (2026-07-14, ericgitonga) — Capture time-on-form for submission metrics (closes #18)
- 47905ae0 (2026-07-14, ericgitonga) — Track metrics scripts; keep extras/ data gitignored
- 44a5f747 (2026-07-14, ericgitonga) — Fix Render logs API format in pull_render_logs.py
- 75947c22 (2026-07-15, ericgitonga) — Add privacy statement to intake form (closes #19)
- d8c6a854 (2026-07-15, ericgitonga) — Introduce semantic versioning and CHANGELOG (closes #20)
- c6b6e2d3 (2026-07-15, ericgitonga) — Add inference-discipline rule to SKILL.md (closes #21)
- dd9d5859 (2026-07-15, ericgitonga) — Second security audit: fix log injection and 4 other findings (closes #22)
- 0668958c (2026-07-15, ericgitonga) — Add warm encouragement banner to Document Uploads section (closes #23)
- 7ed72dd2 (2026-07-15, ericgitonga) — Require CV/business-profile or fallback fields before submission (closes #24)
- 6fed544b (2026-07-15, ericgitonga) — Reword organisation-stage labels away from US venture-funding jargon (closes #25)
- 973d6a35 (2026-07-17, ericgitonga) — Simplify loading page to spinner + tagline, remove countdown timer (closes #26)
- 2b090c0c (2026-07-17, ericgitonga) — Increase loading page tagline font size for legibility (closes #27)
- 61740ae1 (2026-07-17, ericgitonga) — Increase loading page brand text to match tagline size (refs #26)
- 402b8ebb (2026-07-17, ericgitonga) — Require acronyms to be spelled out on first use in generated plans
- 0cfd87a2 (2026-07-17, ericgitonga) — Backfill tags/releases, then scrub client names from every file, tracked or not (closes #29)
- d3a33232 (2026-07-17, ericgitonga) — Extract shared report_builder.py + generate_plan.py CLI to replace per-client scripts (closes #31)
- 3064a9be (2026-07-17, ericgitonga) — Move design_process.pdf and security.pdf output from Clients/ to extras/ (closes #32)

- bf5452a6 (2026-07-17, ericgitonga) — Track generate_design_pdf.py — it never held client data (closes #33)
- 9b7f5de9 (2026-07-28, ericgitonga) — Note the new career-transition landing page repo
- d7a42631 (2026-07-29, ericgitonga) — Align intro and privacy copy with the landing page's brand voice (closes #34)
- 185abe11 (2026-07-29, ericgitonga) — Align loading page copy with the landing page's brand voice (closes #35)
- 2eb4100e (2026-07-29, ericgitonga) — Note that the versioning policy also covers the loading page (closes #36)
- 0308959c (2026-07-30, ericgitonga) — Add payment and platform-unification spike docs (refs #37, refs #38)
- 99956c9f (2026-07-30, ericgitonga) — Track spike/deep-dive PDFs instead of their source markdown (refs #38)
- 3daaa115 (2026-07-30, ericgitonga) — Resolve Option A vs B with real submission-size data (refs #38)
- cf26bd85 (2026-07-30, ericgitonga) — Lower upload cap from 10MB to 4MB (Option A, refs #38)
- a1927905 (2026-07-30, ericgitonga) — v0.21.0: move static assets to public/ for Vercel compatibility
- c59b7413 (2026-07-30, ericgitonga) — v0.21.1: add vercel.json
- ce8a0d9f (2026-07-30, ericgitonga) — v0.21.2: make Flask-Limiter's storage backend configurable
- 5648a48d (2026-07-30, ericgitonga) — Resolve Upstash connection-string question; ignore local tooling files (refs #38)
- 8ec374f5 (2026-07-30, ericgitonga) — v0.21.3: fix vercel.json framework detection and add .vercelignore
- 27c762c6 (2026-07-30, ericgitonga) — v0.21.4: always zip the emailed attachment, even with zero uploads
- 0e2e99f5 (2026-07-30, ericgitonga) — v0.21.5: document Alex Mercer as the standard form-testing fixture
- a9d8ac87 (2026-07-30, ericgitonga) — v0.21.6: document the issue-comment convention this session adopts
- 7015e50e (2026-07-30, ericgitonga) — v0.21.7: untrack SKILL.md
- cdbb7999 (2026-07-30, ericgitonga) — v0.21.8: remove Render-only files and references

Issues (40 tickets, all comments in full)

Issue #1: Initial onboarding form on Hugging Face Spaces

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Build the first version of the career transition intake form using Gradio on Hugging Face Spaces.

Commits: fa58da1, 6310f7a, 7e4c88c, 539ba22

Work done:

- Initial Gradio form with 9 intake sections
- HF Spaces deploy workflow
- Fixed invalid colorTo value, HfFolder import error, and css/theme API change

Outcome: Abandoned — HF Spaces required a paid plan after ZeroGPU lock.

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #2: Platform migration: HF Spaces → Cloud Run → Render (Gradio)

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Evaluate and migrate hosting from HF Spaces through Google Cloud Run to Render, retaining Gradio as the UI framework.

Commits: 852f9e4, 963c651, 08fab6c, a5aa717, 6eb4473, ffc6cf9, 8f9722b, 29d314a, 5e4fb06

Work done:

- Google Cloud Run spike (Dockerfile + GitHub Actions) — aborted due to operational complexity
- Render deployment with Gradio: fixed blank page (Gradio 5→4 downgrade), Python 3.14/audioop crash, boolean schema bug, starlette API change, localhost health-check block

Outcome: Render stable, but Gradio monkey-patches represented unacceptable technical debt → Flask rewrite.

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #3: Rewrite application in Flask + Bootstrap 5

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Replace Gradio entirely with a Flask application and Bootstrap 5 dark-theme form.

Commits: de53d0a, 9ab0104, 82383df, 5306675

Work done:

- Flask app with GET / and POST /submit routes
- ReportLab PDF generation with branded cover block and 9 intake sections
- Bootstrap 5.3 dark theme with navy/teal/gold palette
- Multi-file upload field; headshot field removed
- AJAX fetch submission with PDF blob download and X-Email-Status header

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #4: Switch email delivery from SMTP to Resend API

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Replace smtpplib SMTP delivery (blocked on Render port 587) with the Resend transactional email API over HTTPS.

Commits: c4a29af, 9e00762

Work done:

- Initial SMTP attempt with Gmail App Password — blocked by Render
- Migration to Resend SDK; API key stored as Render environment variable
- Email carries intake PDF and all uploaded documents as attachments

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #5: Client data folder structure and gitignore policy

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Establish Clients/ as the canonical folder for all client data and permanently gitignore it.

Commits: 36a56c2, 22517ab, 1934952

Work done:

- Move all client data under Clients/; update .gitignore
- Add extras/ folder for internal reference documents
- Untrack private PDF generators and report files (security.pdf, generate_security_pdf.py)

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #6: Documentation: hosting guide and design process report

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Produce hosting.pdf (setup guide) and design_process.pdf (project narrative) as ReportLab-generated PDFs.

Commits: 2d7a132, 7b18c00, fdd986a, caf552d, b9bc80e

Work done:

- generate_hosting_pdf.py: illustrated setup guide covering Render, Resend, and maintenance
- generate_design_pdf.py: full design process report (10 sections, platform journey, security architecture, cost analysis)

- generate_selfhosting_pdf.py: self-hosting alternative guide
- Comprehensive docstrings added to all generator functions and app.py

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #7: Security audit and hardening: all 15 findings resolved

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Conduct a manual OWASP-aligned code review and remediate all identified vulnerabilities across three phases.

Commits: 980f026, 5d4301d, 8485fc4, c8a2ae6

Findings and remediations:

- S-01 CSRF protection (Flask-WTF) — HIGH
- S-02 Upload size cap 10 MB — HIGH
- S-03 File extension whitelist — HIGH
- S-04 SRI hashes on CDN resources — HIGH
- S-05 Rate limiting (Flask-Limiter) — MEDIUM
- S-06 Generic error messages to client — MEDIUM
- S-07 Temp file cleanup in finally block — MEDIUM
- S-08 HTTP security headers — MEDIUM
- S-09 SECRET_KEY from environment — MEDIUM
- S-10 Pinned dependencies — MEDIUM
- S-11 Input length limits (_clip) — MEDIUM
- S-12 Control-character stripping (_sanitize) — LOW
- S-13 Random temp filenames — LOW
- S-14 Extension from whitelist only — LOW
- S-15 Structured submission logging — LOW

JavaScript extracted to static/form.js to enable strict script-src CSP.

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #8: Round 1 UX feedback: form flow, copy, and error handling

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Implement feedback from first round of user testing covering form flow, question wording, and error messaging.

Commits: 2ba1e8c, d6a1191

Work done:

- Form flow improvements: section ordering and accordion defaults
- New question added based on tester input
- UX copy refinements across all sections
- Error handling improvements in form.js: plain-English messages for all failure states

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #9: Cold-start loading page to replace Render's server-log screen

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Add a branded Render Static Site as the canonical client entry point, replacing Render's server-log cold-start screen.

Commits: 4d20731, 5736c2f, ffc89d1, 8bf4c0a, 12650ad

Work done:

- loading/index.html: dark-themed branded waiting screen with spinner, progress bar, and 40s countdown
- Polls /_health every 3 seconds; redirects automatically on 200 response
- /_health endpoint with CORS headers scoped to loading site origin (LOADING_SITE_ORIGIN env var)
- Render Static Site service added to render.yaml
- SKILL.md updated with user-facing behaviour rules (plain-English errors, loading page as only client URL)
- design_process.pdf updated to document the cold-start UX decision

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #10: Name intake files by client slug and bundle uploads into a ZIP

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Replace timestamp-based filenames with a client slug (e.g. AlexM) and bundle all uploads into a single ZIP attachment.

Commits: 61bc1a2

Work done:

- `_client_slug()` derives readable identifier from full name (FirstNameInitials)
- `_build_zip()` bundles intake PDF + all uploads into a single ZIP archive
- Email attaches one ZIP (AlexM.zip) instead of multiple loose files
- Bare intake PDF attached directly when no uploads are present

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #11: SKILL.md: specify expertise and complete workflow documentation

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Update SKILL.md with consultant expertise framing, complete onboarding questions, and end-to-end workflow documentation.

Commits: 7773cd5, 693c595

Work done:

- Remove Gradio traces; update SKILL.md to reflect hosted Flask form workflow
- Specify expert HR and Career Coach practitioner framing
- Document all onboarding questions with rationale
- Document 'After intake' mapping table (answers → plan sections)
- Document information to extract from CV, JD, and learning plan

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #12: Round 3: client segmentation, CV fallback, and work-style fields

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Expand the intake form to capture richer client context based on consultant interview.

Commits: 7c4ec23

New fields added:

- Section 1: client type radio, current domain, languages, LinkedIn URL

- Section 3: conditional role title field (shown only for Very clear / Clear target clarity)
- Section 4: work style preference, travel willingness, people management preference
- Section 9: existing portfolio / work evidence textarea
- Section 10: LinkedIn URL field alongside file upload; CV fallback block (toggle-revealed, captures title, industry, years of experience, certifications, key skills)

Also updated: SKILL.md intake questions and extraction notes; generate_design_pdf.py with new design decisions and Round 3 effort row (total: 63 hrs)

1 comment(s):

ericgitonga commented on 2026-07-14:

Closed retrospectively — work completed prior to issue tracking being adopted.

Issue #13: Fix CV fallback position and add free-form 'anything else' field

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Two items from screenshot review (2026-07-14):

Bug — CV fallback block appears under LinkedIn upload, not under CV field When the fallback toggle is clicked, the background fields render below the LinkedIn export upload rather than directly beneath the CV file input. Root cause: the col-12 fallback div sits after the LinkedIn file col-md-6 in the Bootstrap grid. Fix: move the fallback block inside the CV's col-md-6 and convert its internal layout to single-column (stacked fields) to fit the narrower context.

Enhancement — free-form 'anything else' field Add an optional open-text area at the end of Section 9 where clients can capture anything pertinent to their search that no other question covers.

Issue #14: Replace portfolio textarea with dynamic link fields

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Replace the free-text 'existing portfolio' textarea in Section 9 with a structured interaction:

1. Yes/No radio: 'Do you have existing portfolio work to share?'
2. If Yes: reveal a list of URL input fields
3. '+ Add another link' button appends new fields dynamically
4. Each field has an × remove button
5. Links submitted as a list and displayed in the intake PDF

Replaces existing_portfolio (textarea) with portfolio_has_work (radio) + portfolio_links (list of URLs).

Issue #15: Update SKILL.md: visit portfolio links during plan generation

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Claude has WebFetch capability and can visit public URLs as part of plan generation.

Update SKILL.md to instruct the plan generator to:

1. Check portfolio_has_work from the intake PDF
2. Fetch each portfolio link before drafting Section 8
3. Extract relevant context (tech stack, domain, quality signal, recency)
4. Note inaccessible links (auth-gated, 404) and proceed with what's available

No code changes — SKILL.md workflow update only.

Issue #16: Add client-type conditional logic and improve org stage descriptions

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Context

User feedback from beta testers via WhatsApp screenshots captured 2026-07-14:

1. Employment question — entrepreneurs who have never been employed should not see employment-status questions at all; Eric confirmed: "if you have never been employed, you shouldn't have to see it."
2. Org stage descriptions — users didn't understand the startup stage labels; Eric confirmed: "This is the size/stage of a startup company you would like to work for/with. I will revise to make them better explained."

Changes required

1. Client-type conditional logic (form.js + index.html)

- When `client_type` is Entrepreneur or business owner or Freelancer or independent consultant: hide the employment-status block in Section 4 (Constraints & Capacity) and show an entrepreneur-specific alternative question instead
- When `client_type` is Job-seeker / employee in transition or Other: show the existing employment questions
- The "If employed — which best describes your situation?" follow-up already shows/hides correctly but should also be scoped to the job-seeker path

2. Organisation stage descriptions (index.html)

Replace the bare stage labels with richer, self-explaining descriptions so users don't need to guess what each funding/size stage means.

3. SKILL.md + design_process.pdf

Update both documents to reflect the conditional logic flow.

Issue #17: Show clarification text box for any 'Other' response across the form

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

General rule: whenever a question has an Other option and the user selects it, reveal a text input asking them to clarify what they mean.

Currently only one question has this option:

- Section 1 — client_type: "Which best describes you?" → Other

Changes needed:

- index.html: add a d-none text input beneath the client_type radio group, revealed when Other is selected
- form.js: add change listener on client_type radios to show/hide the clarification input; apply the same pattern to any future Other options
- app.py: collect client_type_other in the data dict and include it in the PDF Section 1
- SKILL.md: note the clarification field

Future: if new Other options are added to any question, the same pattern must be applied.

Issue #18: Add submission metrics: Render log extraction + onboard_metrics.xlsx

Opened by ericgitonga on 2026-07-14 — CLOSED — closed 2026-07-14

Add a local-only metrics system to track intake and plan-generation activity.

Sheet 1 — intake (populated by extras/pull_render_logs.py):

- Submission date/time, client name, email, target domain, uploads count
- Time client spent on form (captured as hidden field, logged server-side)
- Server processing time (derived from Render log timestamps)
- Email status, attachment name

Sheet 2 — processed (populated by extras/log_processed.py):

- Client name, processing start/end, duration (min), output file, notes

Tracked changes:

- index.html — hidden time_on_form_seconds field
- form.js — compute elapsed time from page load to submit
- app.py — include time_on_form_seconds in the submission received log line

Local-only (gitignored via extras/):

- extras/pull_render_logs.py — Render API log fetcher + Excel writer

- extras/log_processed.py — manual plan-generation timer logger
- extras/onboard_metrics.xlsx — the output file

Documented in generate_design_pdf.py (also local).

Issue #19: Add privacy statement to intake form

Opened by ericgitonga on 2026-07-15 — CLOSED — closed 2026-07-15

Add a plain-print privacy statement at the bottom of the intake form (above the submit button) clarifying that submitted information/documents are used only to generate the client's Career Transition Plan and are not shared with any third party.

Text to add:

> Your privacy > The information and documents you share in this form are used only to prepare your personalised Career Transition Plan. Nothing you submit is shared with, sold to, or used by any third party, and it is not used for any purpose beyond generating your plan. Your files are handled confidentially and deleted from our servers immediately after your submission is processed.

Also update SKILL.md and design_process.pdf to document this addition.

Issue #20: Introduce semantic versioning and CHANGELOG

Opened by ericgitonga on 2026-07-15 — CLOSED — closed 2026-07-15

Retroactively apply semantic versioning (major.minor.patch) to the project's full commit history, staying pre-1.0 (major=0) since this is an internal tool not yet declared stable.

Plan:

- Group the 53 existing commits into ~24 logical versions (collapsing multi-commit sagas like the Gradio/HF Spaces/Cloud Run/Render hosting migration into single version bumps)
 - Add a CHANGELOG.md (Keep a Changelog style) documenting every version from v0.1.0 (initial commit) to current
 - Add a VERSION file holding the current version string
 - Create annotated git tags (v0.1.0 ... v0.15.0) on the historical commits marking each version boundary
 - Update SKILL.md with the current version and a versioning policy: MINOR for new features/user-facing changes, PATCH for fixes/docs/refactors/hygiene, MAJOR reserved for a future stable 1.0 declaration or breaking changes after that
 - Document the versioning policy as a design decision in generate_design_pdf.py / design_process.pdf
-

Issue #21: Add inference-discipline rule to SKILL.md

Opened by ericgitonga on 2026-07-15 — CLOSED — closed 2026-07-15

Add an explicit rule to SKILL.md: every claim in the generated plan must trace back to something the client explicitly wrote/said/uploaded. No assumptions about anything unclear or unstated. Explicit connections between two stated facts (e.g. a stated preference + a stated skill) are encouraged when they support the client's own stated wishes; connections that require inventing an unstated preference or fact are not allowed.

Issue #22: Second security audit: fix log injection, XLSX formula injection, and 3 other findings

Opened by ericgitonga on 2026-07-15 — CLOSED — closed 2026-07-15

Follow-up security audit covering everything changed since the last full audit (5d4301d, S-01-S-15). Findings to fix, continuing the S-numbering:

- S-16 (HIGH) — Log injection (CWE-117): `full_name`, `client_email`, `target_domain`, `time_on_form_seconds` are logged via `app.logger.info()` without `_sanitize()`, unlike the email-body path. A client can embed a newline to forge a fake 'Submission received:' log line that `extras/pull_render_logs.py`'s regex parser writes into `onboard_metrics.xlsx` as a real submission. Confirmed with a working reproduction.
- S-17 (HIGH) — Formula injection (CWE-1236): client-controlled fields flow unescaped into `openpyxl` cells in `pull_render_logs.py`. A `full_name` of `=HYPERLINK( . . . )` becomes a live formula when the metrics workbook is opened in Excel.
- S-18 (MEDIUM) — Real Render service ID (`srv-d97lh5etrd3s7399aqc0`) hardcoded in `extras/pull_render_logs.py`, a tracked file in this public repo.
- S-19 (LOW) — `_client_slug()` only strips control characters, not path separators; an adversarial `full_name` can embed `../` sequences into ZIP archive entry names (`arcname`). Not exploitable via Python's own `zipfile.extractall()`, which already neutralizes this, but worth restricting as defense in depth.
- Hygiene — stray, superseded `loading.html` at repo root (pre-dates the move to `loading/index.html`), still tracked and public; delete it.

Plan: sanitize all logged fields, neutralize formula-trigger prefixes before writing spreadsheet cells, replace the real service ID with a placeholder, restrict the client slug to safe filename characters, remove the stray file. Update `app.py`'s S-XX docstring, SKILL.md, and the private `generate_security_pdf.py` / `Clients/security.pdf` audit record.

Issue #23: Add warm encouragement banner to Document Uploads section

Opened by ericgitonga on 2026-07-15 — CLOSED — closed 2026-07-15

Two recent intakes were submitted with no CV upload and no fallback background fields filled in, forcing a gap-note follow-up instead of a full plan.

Add a non-blocking, warmly-worded banner in Section 10 (Document Uploads) encouraging clients to fill in what they can — framed as "the more we have, the better the plan," not as a warning or requirement. Static copy, no JS validation gating.

Issue #24: Require CV/business-profile or fallback fields before submission, routed by client type

Opened by ericgitonga on 2026-07-15 — CLOSED — closed 2026-07-15

Follow-up to #23. Both prior thin-data intakes were job-seeker/freelancer types, not entrepreneurs — the CV field's "(optional for entrepreneurs)" label was never actually enforced for anyone, so the non-blocking banner alone won't guarantee future submissions have enough material.

Changes:

- Reuse Section 1's existing `client_type` answer (not a new question) to decide whether Section 10 asks for a CV or a Business Profile/Pitch Deck — same grouping already used for the Section 4 employment-vs-business-status split.
 - Make submission require either that upload or at least one CV-fallback field, enforced both client-side (`form.js`) and server-side (`app.py`), with plain-English error messaging per `SKILL.md`'s error-handling rules.
 - Reword the field label/hint dynamically based on client type.
-

Issue #25: Replace US venture-funding jargon (seed/Series A/Series B) in organisation-stage field with geography-neutral language

Opened by ericgitonga on 2026-07-15 — CLOSED — closed 2026-07-15

The Preferred organisation stage checkboxes on the intake form (`templates/index.html`) label the two smaller tiers as:

- "Early-stage startup (seed / Series A)"
- "Growth-stage startup (Series B+)"

This is US venture-capital funding-round terminology. The primary client base is Kenya-based, where most employers a career-transition client targets are not VC-funded in the Silicon Valley sense — assuming seed/Series A/Series B as the framing for "early" vs "growth" doesn't reflect the local business landscape.

Fix: reword the tier labels to describe team size and organisational maturity only (which the existing muted description text already does), dropping the funding-round parenthetical entirely. Keep the same 5 tiers (early / growth / established / large / no-preference) and the same `org_stage` form-field values structurally compatible — only the microcopy changes.

No new org types added in this pass (scoped down from a broader Kenyan-org-type expansion after discussion).

Issue #26: Simplify loading page tagline and remove timer

Opened by ericgitonga on 2026-07-17 — CLOSED — closed 2026-07-17

Replace 'Application is starting up' / 'Estimated wait: Ns' countdown with a simple tagline in the style of PDF Drive's loading screen (e.g. 'Preparing your form...'), and remove the visible countdown timer text. Progress bar can remain as a visual indicator without the numeric ETA.

Issue #27: Increase loading page tagline font size

Opened by ericgitonga on 2026-07-17 — CLOSED — closed 2026-07-17

The 'Preparing your form...' tagline on the loading page (loading/index.html) was too small at .8rem. Increase to ~24px (1.5rem) for better legibility.

1 comment(s):

ericgitonga commented on 2026-07-17:

Duplicate of #26 — this tweak should have been folded into that issue rather than filed separately.

Issue #28: Spell out acronyms in full on first use in generated plans

Opened by ericgitonga on 2026-07-17 — CLOSED — closed 2026-07-17

Generated career transition plans currently drop acronyms (e.g. IoD, COI) into the document without ever spelling out what they stand for, leaving the client to guess. Every acronym should be written in full the first time it appears in a plan, with the acronym in parentheses, then the acronym alone can be used for the rest of the document.

Scope:

- Fix an existing client plan (generate_plan.py + regenerate the affected transition_plan.pdf)
 - Update SKILL.md with this as a standing rule for all future plans
 - Update generate_design_pdf.py and regenerate design_process.pdf to document the rule
-

Issue #29: Backfill missing git tags and add GitHub Releases to the versioning workflow

Opened by ericgitonga on 2026-07-17 — CLOSED — closed 2026-07-17

CHANGELOG.md claims tags v0.18.0, v0.19.0, v0.19.1, and v0.19.2 exist, but they were never actually created/pushed — only v0.1.0 through v0.17.0 exist as real git tags. Additionally, no GitHub Release has ever been published for any version.

Scope:

- Create and push the 4 missing tags at the correct commits (per CHANGELOG's closes-# references)
 - Retroactively create a GitHub Release for every existing tag (v0.1.0 through v0.19.2), with release notes derived from CHANGELOG.md and a Full-Changelog compare link between consecutive versions, similar to GitHub's own auto-generated release format
 - Update SKILL.md and generate_design_pdf.py/design_process.pdf so that, going forward, every version bump also creates + pushes a git tag and publishes a matching GitHub Release before/alongside the commit
-

Issue #30: Scrub client names from CHANGELOG, SKILL.md, and published GitHub Releases

Opened by ericgitonga on 2026-07-17 — CLOSED — closed 2026-07-17

The v0.29 tag/release backfill (issue #29) generated GitHub Release notes directly from CHANGELOG.md, which included client names in v0.18.0, v0.19.0, and v0.19.2 — now public and permanent record.

Scope:

- Finish scrubbing client names from CHANGELOG.md (referring to incidents generically, e.g. "a client plan", "two prior thin-data intakes")
- Add a standing rule to SKILL.md: never name a client in CHANGELOG.md, a commit message, a git tag, or a GitHub Release — these are public/tracked; generate_design_pdf.py/design_process.pdf remain fine for internal naming since they're gitignored
- Edit the 3 affected published GitHub Releases (v0.18.0, v0.19.0, v0.19.2) to remove client names, matching the corrected CHANGELOG.md wording
- Update generate_design_pdf.py/design_process.pdf to document this decision

1 comment(s):

ericgitonga commented on 2026-07-17:

Duplicate — this scope was already folded into #29's implementation (already staged locally: CHANGELOG v0.19.3 entry, VERSION bump, SKILL.md scrub rule). Continuing there instead of splitting the work.

Issue #31: Extract shared report_builder.py + generate_plan.py CLI to replace per-client scripts

Opened by ericgitonga on 2026-07-17 — CLOSED — closed 2026-07-17

Every client folder currently has its own generate_plan.py (900-1200+ lines), duplicating ~250 lines of identical ReportLab boilerplate (palette, ParagraphStyles, section_header/shaded_box/two_col/semester_card/data_table/footer_canvas helpers) at the top, with only the remainder being actual client content.

Scope:

- Extract the shared boilerplate and per-SKILL.md-section rendering logic into a single root-level, tracked report_builder.py (build_plan(data, output_path) plus one render_*() function per document section)
- Add a thin root-level generate_plan.py CLI: python3 generate_plan.py "<Client Name>" — loads Clients/<Name>/plan_data.py's PLAN dict and calls report_builder.build_plan()
- Define the plan_data.py schema so client folders hold content only (no ReportLab code), with variable-length sections (semester count, archetypes, optional Year-2 track, entrepreneur vs job-seeker framing) falling out of plain list/dict length rather than special-cased code
- Migrate Alex Mercer's existing generate_plan.py into Clients/Alex Mercer/plan_data.py as the round-trip proof, confirming the rendered PDF matches the original in structure and content
- Existing real clients' generate_plan.py scripts are left as-is (already-delivered artifacts) unless a future change asks to migrate them too
- Update SKILL.md (PDF generation approach, workflow steps, files-to-keep-in-client-folder) and generate_design_pdf.py/design_process.pdf to document the new architecture

Issue #32: Move design_process.pdf and security.pdf output from Clients/ to extras/

Opened by ericgitonga on 2026-07-17 — CLOSED — closed 2026-07-17

Both generate_design_pdf.py and generate_security_pdf.py write their output (design_process.pdf, security.pdf) into Clients/, alongside actual client folders — even though both are internal project documents, not client data. This caused confusion locating the file, since extras/ is where every other non-client-specific gitignored project artefact already lives (feedback_01.pdf, costs.pdf, the metrics spreadsheet), and a stale extras/design_process.pdf copy from earlier in the project already existed there.

Scope:

- Update both scripts' OUTPUT_PATH (and docstrings) to write into extras/ instead of Clients/, matching the convention generate_feedback_pdf.py already follows
- Move the existing design_process.pdf and security.pdf files accordingly
- Update SKILL.md's client-data-handling rules to reflect the new location
- Regenerate both PDFs and update generate_design_pdf.py's own Key Design Decisions section to document the move

Issue #33: Track generate_design_pdf.py — it contains no client data

Opened by ericgitonga on 2026-07-17 — CLOSED — closed 2026-07-17

generate_design_pdf.py is currently gitignored under the 'Private report generators' block alongside generate_security_pdf.py, generate_feedback_pdf.py, and generate_selfhosting_pdf.py. Unlike those, it contains no client data or real secrets — just project/process narrative (architecture decisions, effort log, cost analysis), the approved fictitious Alex Mercer example, and generic placeholders (env var names like SECRET_KEY, a 'Jane Doe' CLI usage example). Its output (design_process.pdf) stays gitignored in extras/ as a generated artefact, but the generator script itself has no reason to be untracked.

Scope:

- Remove generate_design_pdf.py from .gitignore's private-generators block
- git add and commit it
- Update SKILL.md's client-data-handling rules to reflect its tracked status
- Document the change in generate_design_pdf.py's own Key Design Decisions section and regenerate design_process.pdf

Issue #34: Align intro and privacy copy with the career-transition landing page's brand voice

Opened by ericgitonga on 2026-07-29 — CLOSED — closed 2026-07-29

The new marketing landing page (ericgitonga/career-transition) has its own distinct voice — confident, second-person, benefit-led ("Your next career isn't a guess. It's a plan.") — that the intake form's copy

predates and doesn't match. A visitor clicking through from the landing page's CTA currently lands on noticeably flatter, more procedural copy.

Reword two spots in templates/index.html to match the landing page's tone (same substantive meaning, no new fields/behaviour):

- The form-header intro (h1 + two lead paragraphs above the accordion).
- The privacy paragraph above the submit button — keep every existing commitment (used only to build the plan, never shared/sold/used by a third party, files deleted from the server immediately after processing), just in the landing page's voice.

Issue #35: Align loading page copy with the career-transition landing page's brand voice

Opened by ericgitonga on 2026-07-29 — CLOSED — closed 2026-07-29

Follow-up to #34 (intro/privacy copy). The cold-start loading page (loading/index.html, deployed as the career-transition-loading Render static site — the actual public entry point clients land on from the landing page's CTA) still says "Career Transition Onboarding", which predates the landing page's branding ("Career Transition Planning").

Update the brand wordmark/title and the status tagline to match, without changing the polling/redirect behaviour.

Issue #36: Note that the versioning policy also covers the loading page

Opened by ericgitonga on 2026-07-29 — CLOSED — closed 2026-07-29

SKILL.md's Versioning section didn't explicitly say the repo's single VERSION/CHANGELOG.md line also covers loading/index.html (deployed separately as the career-transition-loading Render static site) — it read as if only app.py/templates/index.html were in scope. #35's loading-page copy rework (v0.20.4) was in fact versioned this way, but the policy doc didn't say so.

Add an explicit note, and fix the "Current version" line (stale at 0.20.2) plus a quoted tagline string that's now out of date since #34/#35 changed the actual copy.

Issue #37: Spike: collect payment (M-Pesa, Pesapal) on intake form submission

Opened by ericgitonga on 2026-07-30 — OPEN

Explore how to actually collect the advertised intake fee (currently KES 7,500 introductory) at the moment a client submits the intake form, starting with M-Pesa and Pesapal.

Covers: direct M-Pesa Daraja vs. Pesapal vs. other aggregators (onboarding realism for a solo/informal consultant, integration effort, fees), the two-step submit-then-pay-then-confirm flow needed to gate /submit on payment, failure modes (Render cold-start vs. payment webhooks), and logging every transaction to a Google Sheet as a simple ledger.

Spike only — no implementation yet. See extras/spike_payment_integration.md.

1 comment(s):

ericgitonga commented on 2026-07-30:

Spike complete — see [extras/spike_payment_integration.pdf](#) (tracked; the source markdown is kept local, not in git).

Provider recommendation: Pesapal over direct M-Pesa Daraja, despite Daraja's lower transaction fee (<1% vs Pesapal's 1.5-3.5%). Direct Daraja requires personally holding a Paybill/Till — realistically 4-8 weeks of KYC/go-live bureaucracy for a first-timer before a single payment can be collected. Pesapal already holds its own Daraja relationship and lets a solo/informal consultant onboard as a sub-merchant with lighter KYC (ID, KRA PIN, no shortcode application), and its API v3/IPN integration is genuinely simple — roughly 1-2 days of Flask work vs 2-4 days of Daraja code plus the multi-week credential wait. IntaSend was flagged as worth a same-day pricing/onboarding check before committing, since it markets itself as friendlier still to an unregistered solo business.

Integration pattern: /submit needs to become two-step — validate and hold form state server-side first, trigger payment (Pesapal hosted checkout redirect), then only build the PDF/email it once a webhook/GetTransactionStatus call confirms payment, never trusting the browser redirect alone. Flagged a real interaction with Render's free-tier cold start: a payment webhook landing on a cold instance risks being dropped, which argues for either upgrading off the free tier or leaning on query-based status checks as the source of truth once payment actually gets built.

Added mid-spike (at the app owner's request): a Google Sheets transaction ledger, via gspread + a service account + a Render Secret File for the credential — reuses the same "external API + Render-stored secret" pattern the app already has for Resend, estimated well under an hour of added work on top of the payment integration itself. Recommended over a Google Apps Script webhook alternative, which would've traded proper IAM-based access control for a hand-rolled shared-secret check on a publicly-POSTable URL — the wrong tradeoff for something logging financial transactions.

Still open, not part of this spike: whether the consultant is a registered sole proprietorship or fully informal (determines which providers are realistically approvable), and confirming IntaSend's actual fee sheet directly.

Issue #38: Spike: unify landing page + intake app under one domain

Opened by ericgitonga on 2026-07-30 — OPEN

Explore how to make the Vercel-hosted landing page and the Render-hosted intake app (plus its loading/cold-start page) feel like one product to a visitor — no visible onrender.com hop, ideally one hostname throughout.

Covers: DNS subdomain split, a Vercel rewrite proxying to Render, a Cloudflare Worker reverse proxy in front of both, and fully migrating the Flask app onto Vercel — with effort/maintenance/completeness tradeoffs and a ranked recommendation.

Spike only — no implementation yet. See [extras/spike_platform_unification.md](#).

1 comment(s):

ericgitonga commented on 2026-07-30:

Long arc on this one — recapping the real decisions and course-corrections, per SKILL.md's convention.

Direction shift: the initial spike ([extras/spike_platform_unification.pdf](#)) compared four options (DNS subdomain split, Vercel rewrite to Render, a Cloudflare Worker reverse proxy, and a full Flask-to-Vercel migration) and recommended trying the Vercel rewrite first, since the full migration was blocked by two real unknowns: Vercel's request-body limit against this app's 10MB upload requirement,

and Flask-Limiter's in-memory rate-limit storage not working across serverless instances. The app owner chose to lean into the full migration anyway, so a deep dive ([extras/deepdive_flask_to_vercel.pdf](#)) went after both blockers directly.

Resolved a stale claim before trusting it: initial research surfaced a "100MB body limit" figure from a bundled knowledge note, contradicting the 4.5MB figure in older docs. Fetched Vercel's live limits page directly rather than trusting either source blind — confirmed 4.5MB is still the real, current ceiling.

Data beat assumption on Option A vs B: rather than guess whether shrinking the upload cap (Option A, simple) or building a Vercel Blob client-upload workaround (Option B, ~2-3 extra days, needs a second Node runtime since there's no Python-native Blob token-issuing path) was viable, checked the actual retained files from ~12 past client engagements in `Clients/`. Largest real upload on record: ~445KB. Largest plausible combined submission: ~654KB. Both nowhere near even a conservative 4MB cap. Option A confirmed, Option B shelved as a fallback — dropped the migration effort estimate from 4-6 days to 2-3.

Shipped the migration steps (v0.21.0-0.21.2): lowered `MAX_CONTENT_LENGTH` to 4MB, moved `static/form.js` to `public/form.js` with `static_url_path=""` (works identically on Render, matches what Vercel's Flask guide requires), added `vercel.json`, made Flask-Limiter's storage configurable via `RATELIMIT_STORAGE_URI`.

Provisioning surprised the plan, for the better: the deep dive couldn't confirm from Vercel's own docs whether the Upstash Marketplace integration exposes a real `redis:///rediss://` TCP string (Flask-Limiter's actual requirement) or only the REST vars Vercel's SDK docs describe. Actually provisioning it (vercel integration add upstash/upstash-kv against a new career-transition-intake Vercel project) settled it directly: it injects `REDIS_URL` using `rediss://`, exactly what was needed. Copied into `RATELIMIT_STORAGE_URI`.

Two real bugs found only by actually deploying, not by reading docs:

1. An explicit functions: `{"app.py": {maxDuration: 60}}` block in `vercel.json` silently disabled Vercel's zero-config Flask auto-detection — the build "succeeded" in 25ms but produced no Serverless Function at all, so every route 404'd. Fixed with `"framework": "flask"` instead (v0.21.3).
2. `vercel deploy --dry --json`, run deliberately before any real deploy, showed `Clients/` (8.3MB of real client CVs and personal documents), `extras/`, `.agents/`, `.claude/`, and the private report generators would all have been uploaded — `vercel deploy` does not respect `.gitignore`. Added `.vercelignore` and re-verified the exclusion took effect before deploying for real (v0.21.3).

Full smoke test, not just a health check: verified the preview end to end using Alex Mercer's actual original intake data (form page, static asset, `/_health`, and a real multi-file POST `/submit` — CV+JD+LinkedIn, PDF generation, ZIP bundling, Resend call, and the Redis-backed rate limiter genuinely returning 429 after 5 requests/minute). Confirmed via `vercel logs` that all 3 uploads were received server-side (uploads=3), not just that the request returned 200.

Incidental, unrelated finding worth flagging here for the record: while investigating an unfamiliar file during this work, an old `RENDER_API_KEY` sitting in a local, untracked `.claude/settings.local.json` got printed to a session transcript. Confirmed via `git log --all` it was never committed. The app owner rotated the key; `.claude/`, `.agents/`, and `skills-lock.json` are now gitignored so nothing like it can be committed going forward.

Current state: Vercel preview deployment is fully verified but not yet production. Domain cutover and Render decommissioning are next.

Issue #39: Always zip the consultant's emailed attachment, even with zero uploads

Opened by ericgitonga on 2026-07-30 — CLOSED — closed 2026-07-30

Found while smoke-testing the Vercel migration preview (#38) with a full real-world submission.

Before this fix, `\submit\` only built a ZIP for the consultant's emailed attachment when the client uploaded at least one file — with zero uploads, the email attachment was the bare intake PDF instead. This meant the consultant received two different attachment formats depending on client behavior.

Fix: always call `\_build_zip()\`, even with an empty uploads list (it already handles that fine — the intake PDF alone still gets zipped). The consultant now always receives a `\.zip\` from every submission.

1 comment(s):

ericgitonga commented on 2026-07-30:

Found this by accident while smoke-testing #38's Vercel migration with a full real-world submission (Alex Mercer's actual original intake data) — the server log showed `attachment=AlexM.zip` for a submission with uploads, which prompted checking what a zero-upload submission would log, and confirmed the inconsistency. Unrelated to the migration itself; would have reproduced identically on Render.

Fixed in v0.21.4 by removing the `if uploads: zip / else: plain PDF` branch entirely — `\_build_zip()` already handled an empty uploads list correctly, so the fix was just always calling it.

Separately, and not fixed here since it wasn't part of this ask: the browser-facing download (what the client sees after submitting) always streams the plain intake PDF via `send_file(pdf_path, ...)`, regardless of whether a ZIP was built — only the consultant's emailed copy uses the ZIP. That's a distinct, still-open discrepancy if the client-facing download ever needs to include their own uploads too.

Issue #40: Decommission Render: remove render.yaml, loading page, and Render-only references

Opened by ericgitonga on 2026-07-30 — CLOSED — closed 2026-07-30

Follow-up to #38, after the Vercel migration was verified and cut over: both Render services (`career-transition-intake`, `career-transition-loading`) have been deleted, and the landing page's CTA now points directly at `https://career-transition-intake.vercel.app` in production.

This tracks cleaning up everything that only existed to serve the now-deleted Render setup:

- `render.yaml` — described services that no longer exist
- `loading/index.html` and its Render static-site service — existed solely to mask Render's free-tier cold start, which Vercel's Fluid Compute doesn't have
- `app.py`'s `/_health` CORS headers — only needed for the loading page's cross-origin polling
- Stale "share the loading page URL" / Render-specific instructions in `SKILL.md` and `README.md`

1 comment(s):

ericgitonga commented on 2026-07-30:

Shipped in v0.21.8. One judgment call worth recording: kept `\/_health\` itself rather than removing it outright — only stripped its CORS headers, which were the loading-page-specific part. A bare liveness probe still has ongoing value for monitoring/debugging independent of why it was originally added (I used it myself while smoke-testing #38).

Also flagged, not fixed here: \hosting.pdf\ (referenced from README.md) still describes the old Render setup and hasn't been regenerated for Vercel — noted in the README as a known gap rather than silently left wrong.

Releases

47 releases:

v0.1.0 — 2026-07-15

What's Changed

Added

- Initial onboarding intake form and Hugging Face deploy workflow.

v0.1.1 — 2026-07-15

What's Changed

Changed

- Client data moved under Clients/; .gitignore and SKILL.md updated; hosting.md and generate_hosting_pdf.py added to .gitignore.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.1.0...v0.1.1>

v0.2.0 — 2026-07-15

What's Changed

Changed

- Hosting migrated from Hugging Face Spaces through Google Cloud Run to Render; Gradio replaced entirely with Flask after repeated platform-specific crashes (dependency pinning, a localhost-check bypass, boolean-schema patches, and more along the way).

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.1.1...v0.2.0>

v0.3.0 — 2026-07-15

What's Changed

Changed

- Email delivery switched from SMTP to the Resend API, fixing an SMTP worker timeout; App Password setup guide added as an SMTP fallback reference.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.2.0...v0.3.0>

v0.4.0 — 2026-07-15

What's Changed

Changed

- Form UI refresh: dark theme applied, header text updated, multi-file upload field added; Email Settings section and headshot upload field removed.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.3.0...v0.4.0>

v0.4.1 — 2026-07-15

What's Changed

Added

- Hosting guide (`hosting.pdf`) added to the repo; docs updated for the Flask + Resend stack; docstrings added throughout `app.py` and `generate_hosting_pdf.py`.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.4.0...v0.4.1>

v0.5.0 — 2026-07-15

What's Changed

Added

- Design-process report generator (`generate_design_pdf.py`) producing an internal record of the project's design decisions.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.4.1...v0.5.0>

v0.5.1 — 2026-07-15

What's Changed

Removed

- Leftover Gradio references cleaned up; `SKILL.md` updated for the hosted Flask form.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.5.0...v0.5.1>

v0.6.0 — 2026-07-15

What's Changed

Added

- Self-hosting guide generator (`generate_selfhosting_pdf.py`); docstrings added throughout `generate_design_pdf.py`.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.5.1...v0.6.0>

v0.7.0 — 2026-07-15

What's Changed

Security

- All 15 findings from the internal security audit resolved (CSRF, rate limiting, upload validation, security headers, etc.); `SECRET_KEY` auto-generated locally when absent from the environment; private audit script gitignored.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.6.0...v0.7.0>

v0.7.1 — 2026-07-15

What's Changed

Changed

- Design process report updated with security content; added `extras/` folder for local-only tooling.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.7.0...v0.7.1>

v0.8.0 — 2026-07-15

What's Changed

Changed

- Round 1 client feedback applied: form flow adjustments, a new question, UX copy improvements, and better error handling.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.7.1...v0.8.0>

v0.9.0 — 2026-07-15

What's Changed

Added

- Clean branded loading screen replacing Render's default cold-start page; moved to its own static subdirectory.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.8.0...v0.9.0>

v0.9.1 — 2026-07-15

What's Changed

Changed

- User-facing message wording edits; user-facing behaviour rules documented in SKILL.md; design_process.pdf updated with loading-screen UX notes; cold-start feedback captured; private PDF generators and report files untracked from git.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.9.0...v0.9.1>

v0.10.0 — 2026-07-15

What's Changed

Changed

- Intake files now named by client slug; uploads bundled into a single ZIP attachment.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.9.1...v0.10.0>

v0.10.1 — 2026-07-15

What's Changed

Changed

- SKILL.md wording clarified.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.10.0...v0.10.1>

v0.11.0 — 2026-07-15

What's Changed

Added

- Client segmentation, CV fallback block, and work-style / travel / management-preference fields.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.10.1...v0.11.0>

v0.11.1 — 2026-07-15

What's Changed

Fixed

- CV fallback block position corrected; added a free-form notes field. (closes #13)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.11.0...v0.11.1>

v0.12.0 — 2026-07-15

What's Changed

Changed

- Portfolio textarea replaced with dynamic add/remove link fields. (closes #14)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.11.1...v0.12.0>

v0.12.1 — 2026-07-15

What's Changed

Changed

- Plan generator instructed to fetch each portfolio link via WebFetch before drafting the portfolio section. (closes #15)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.12.0...v0.12.1>

v0.13.0 — 2026-07-15

What's Changed

Added

- Client-type conditional logic (job-seeker / entrepreneur / freelancer) driving employment vs. business-status questions; improved organisation-stage descriptions. (closes #16)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.12.1...v0.13.0>

v0.13.1 — 2026-07-15

What's Changed

Added

- Free-text clarification field shown when "Other" is selected for client type. (closes #17)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.13.0...v0.13.1>

v0.14.0 — 2026-07-15

What's Changed

Added

- Submission metrics system: hidden time-on-form field captured on submit; local scripts to pull Render logs and log processed-plan turnaround into a tracked spreadsheet. (closes #18)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.13.1...v0.14.0>

v0.15.0 — 2026-07-15

What's Changed

Added

- Privacy statement on the intake form clarifying submitted data is used only to generate the client's plan and is not shared with any third party. (closes #19)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.14.0...v0.15.0>

v0.16.0 — 2026-07-15

What's Changed

Added

- Semantic versioning introduced: VERSION file, this changelog, retroactive version tags across the full commit history, and a versioning policy documented in SKILL.md. (closes #20)
- Current version displayed at the bottom of the intake form's header, read live from VERSION via a new app_version template variable.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.15.0...v0.16.0>

v0.16.1 — 2026-07-15

What's Changed

Changed

- Added an "Inference Discipline" section to SKILL.md: every claim in the generated plan must trace back to something the client explicitly stated; connections between two stated facts are encouraged, but no assumption may be invented to fill an unstated gap. (closes #21)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.16.0...v0.16.1>

v0.17.0 — 2026-07-15

What's Changed

Security

- Second security audit, covering everything changed since the S-01-S-15 hardening pass: S-16 log-line forgery / field-hijack in the submission log (fixed via sanitizing and percent-encoding logged fields, `_log_field()`); S-17 spreadsheet formula injection in `onboard_metrics.xlsx` (fixed via a leading-apostrophe guard, `_safe_cell()`, in both extras/ scripts); S-18 a real Render service ID hardcoded in a public, tracked script (replaced with a placeholder); S-19 client slug not restricted to safe filename characters, allowing path-like sequences into ZIP archive entry names (fixed by restricting `_client_slug()` to letters/digits/hyphen/underscore). Also removed a stale, superseded `loading.html` left at the repo root. (closes #22)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.16.1...v0.17.0>

v0.18.0 — 2026-07-17

What's Changed

Added

- Static, non-blocking encouragement banner in Section 10 (Document Uploads), inviting clients to fill in a CV, target job description, or the CV fallback fields — framed positively ("the more we have, the better the plan"), never as a warning or a requirement. No validation gating was added; submission is unaffected. Added after two prior intakes arrived with neither a CV nor the fallback fields filled in, which meant no meaningful plan could be generated and a manual gap-note follow-up was needed instead. (closes #23)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.17.0...v0.18.0>

v0.19.0 — 2026-07-17

What's Changed

Changed

- Section 10's non-blocking encouragement banner (0.18.0 / #23) is superseded by an actual requirement: submission now needs either a CV/business-profile upload or at least one CV-fallback field filled in, enforced both client-side (`form.js`) and server-side (`app.py`, HTTP 400 with a plain-English message if bypassed).
- The CV field now reuses Section 1's existing `client_type` answer (no new question) to relabel itself: job-seekers/employees/"Other" see "CV / Résumé"; entrepreneurs and freelancers see "Business Profile / Pitch Deck" — the same grouping already used for Section 4's employment-vs-business split, factored into a shared `_is_entrepreneur_type()` helper in `app.py`.
- Motivated by discovering that both prior thin-data intakes were job-seeker/freelancer types, not entrepreneurs — the old "(optional for entrepreneurs)" label had never actually been enforced for

anyone. (closes #24)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.18.0...v0.19.0>

v0.19.1 — 2026-07-17

What's Changed

Changed

- Reworded the two smaller "Preferred organisation stage" checkbox labels in Section 2 of the intake form: "Early-stage startup (seed / Series A)" → "Early-stage business"; "Growth-stage startup (Series B+)" → "Growth-stage business". The dropped parenthetical was US venture-capital funding-round terminology that doesn't reflect the primary client base (Kenya-based), where most target employers aren't venture-funded in that sense. The five underlying tiers and their team-size/maturity descriptions are unchanged. (closes #25)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.19.0...v0.19.1>

v0.19.2 — 2026-07-17

What's Changed

Changed

- Added a standing content rule to SKILL.md: every acronym or abbreviation a client would not already know must be spelled out in full (acronym in parentheses) the first time it appears anywhere in a generated plan, by reading order, and used as the acronym alone afterwards, including inside tables. Terms already standard in this workflow's vocabulary (CV, NGO, CEO, PDF, currency codes) are exempt. Added a matching quality-checklist item.
- Documented the decision in generate_design_pdf.py's Key Design Decisions section and regenerated design_process.pdf.
- Fixed a client plan generated before this rule existed, where several acronyms were either never expanded or expanded out of reading order. (closes #28)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.19.1...v0.19.2>

v0.19.3 — 2026-07-17

What's Changed

Changed

- Backfilled the git tags CHANGELOG.md had already been claiming existed (v0.18.0, v0.19.0, v0.19.1, v0.19.2) at the commits their issue references point to, and pushed them.
- Published a GitHub Release for every version from v0.1.0 through v0.19.2 (31 total), built from each CHANGELOG.md entry's own subheading and bullets plus a Full-Changelog compare link to the previous tag.

- Reworded three CHANGELOG.md entries (v0.18.0, v0.19.0, v0.19.2) that had named real clients as the motivating example behind a fix, replacing the names with a generic description, and edited the matching GitHub Release bodies to match. Also scrubbed the equivalent mentions from `generate_design_pdf.py`, since a gitignored file is not the same as a safe one.
- SKILL.md now requires a pushed git tag and a published GitHub Release as the last step of every version bump, and explicitly prohibits naming a real client anywhere — tracked or gitignored, published or purely local. Alex Mercer, the fictitious demo client, is the one standing exception and is now the reference example wherever SKILL.md needs to point to a concrete case.
- Added `Clients/Alex Mercer/generate_gap_note.py`, producing a fictitious `gap.pdf` for a hypothetical thin intake (no CV, no JD), so the gap-note pattern referenced in SKILL.md's Document Uploads rule has a nameable example instead of pointing at a real client's folder.
- Documented all three decisions in `generate_design_pdf.py`'s Key Design Decisions section and regenerated `design_process.pdf`. (closes #29)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.19.2...v0.19.3>

v0.20.0 — 2026-07-17

What's Changed

Added

- Extracted a shared `report_builder.py` (root, tracked) holding all ReportLab palette, paragraph styles, layout helpers, and per-section rendering logic that used to be duplicated at the top of every client's own `generate_plan.py` (~250 lines of identical boilerplate each). Added a thin root-level `generate_plan.py` CLI (`python3 generate_plan.py "<Client Name>"`) that loads `Clients/<Name>/plan_data.py`'s PLAN dict and calls `report_builder.build_plan()`.
- Defined the `plan_data.py` schema: fixed-shape sections (2, 3, 5, 8, 9, 11) as structured lists of dicts; freeform sections (1, 4, 6, 7, 10, 12) as a generic blocks list (paragraph / shaded_box / two_col / table) rendered by one shared `render_blocks()` — no per-client ReportLab code anywhere.
- Migrated Alex Mercer's plan to `Clients/Alex Mercer/plan_data.py` as the round-trip proof and removed the old per-client `generate_plan.py`; the rendered PDF is byte-for-byte equivalent in content (17 pages, 5,252 words, confirmed via `pdftotext` diff — the only differences were word-wrap order inside the Monthly Action Tracker table, not missing content).
- Existing real clients' `generate_plan.py` scripts are left as delivered (not retroactively migrated); new clients use the shared engine going forward.
- Updated SKILL.md (PDF generation approach, workflow steps 4–5, files to keep in the client folder) to document the new architecture and updated `generate_design_pdf.py/design_process.pdf` accordingly. (closes #31)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.19.3...v0.20.0>

v0.20.1 — 2026-07-17

What's Changed

Changed

- `generate_design_pdf.py` and `generate_security_pdf.py` now write their output (`design_process.pdf`, `security.pdf`) to `extras/` instead of `Clients/` — both are internal project documents, not client data, and writing them into `Clients/` alongside actual client folders was a source of real confusion when looking for the file. Matches the convention `generate_feedback_pdf.py` already used. Overwrote a stale `extras/design_process.pdf` copy left over from earlier in the project with a freshly regenerated one.
- Updated SKILL.md's client-data-handling rules to reflect the new location and `generate_design_pdf.py`'s own Key Design Decisions section to document the move. (closes #32)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.20.0...v0.20.1>

v0.20.2 — 2026-07-17

What's Changed

Changed

- `generate_design_pdf.py` is now tracked instead of gitignored: a check of its content (every email address, every capitalised two-word phrase, every SECRET_KEY/API_KEY-shaped string) confirmed it holds no client data or real secrets — only project narrative, the approved fictitious Alex Mercer example, a generic "Jane Doe" placeholder, and environment-variable names already documented in tracked SKILL.md and `render.yaml`, never a value. It had been gitignored only by inherited convention (grouped with three other generators that do hold private data), not by its own content. Its output, `extras/design_process.pdf`, stays gitignored as a generated artefact.
- Updated SKILL.md's client-data-handling rules and `generate_design_pdf.py`'s own Key Design Decisions section to document this. (closes #33)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.20.1...v0.20.2>

v0.20.3 — 2026-07-29

What's Changed

Changed

- Reworded the form-header intro and the privacy paragraph in `templates/index.html` to match the new career-transition landing page's brand voice (confident, second-person, benefit-led) instead of the original flatter, procedural copy — same substantive meaning throughout, including every existing privacy commitment (used only to build the plan, never shared/sold/used by a third party, files deleted immediately after processing). (closes #34)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.20.2...v0.20.3>

v0.20.4 — 2026-07-29

What's Changed

Changed

- Reworded the cold-start loading page (`loading/index.html`, the career-transition-loading Render static site — the actual public entry point clients land on from the landing page's CTA) to match the same brand voice as #34: brand wordmark and `<title>` changed from "Career Transition Onboarding" to "Career Transition Planning", and the status tagline reworded ("Preparing your plan..." / "All set — taking you there..."). Polling/redirect behaviour unchanged. (closes #35)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.20.3...v0.20.4>

v0.20.5 — 2026-07-29

What's Changed

Changed

- SKILL.md's Versioning section now explicitly notes that the repo's single `VERSION/CHANGELOG.md` covers `loading/index.html` too, deployed separately as the career-transition-loading Render static site — a copy-only change there still needs its own version bump and dated `CHANGELOG.md` entry, exactly like `app.py/templates/index.html`. Also refreshed the stale "Current version" line and a quoted loading-page tagline string that #34/#35 had already changed in practice. (closes #36)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.20.4...v0.20.5>

v0.21.0 — 2026-07-30

What's Changed

Changed

- Lowered `MAX_CONTENT_LENGTH` from 10 MB to 4 MB. Real submissions to date top out around 450 KB per file, and staying under 4 MB keeps the app deployable as-is on Vercel Functions, whose request body limit is a hard 4.5 MB — the first step of the Flask-to-Vercel migration (see `extras/deepdive_flask_to_vercel.md`). Updated the 413 error message (`static/form.js`) and docstrings/security notes (`app.py`) to match. (refs #38)
- Moved the app's static asset (`form.js`) from `static/` to `public/`, and switched Flask's `static_folder/static_url_path` so it's served at the URL root (`/form.js`) instead of under `/static/`. Required by Vercel's Flask guide, which expects static assets under `public/**` so its CDN can serve them directly instead of through the Function; behaviorally identical on the current Render deployment. (refs #38)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.20.5...v0.21.0>

v0.21.1 — 2026-07-30

What's Changed

Added

- Added `vercel.json`, capping the Flask function's `maxDuration` at 60s — generous for this app's actual synchronous work (PDF build + one Resend call), while bounding worst-case cost well under Hobby's 300s default/max. No relocation needed for `app.py`; Vercel auto-detects it at the project root. Third step of the Flask-to-Vercel migration (see `extras/deepdive_flask_to_vercel.md`). No effect on the current Render deployment, which doesn't read this file. (refs #38)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.21.0...v0.21.1>

v0.21.2 — 2026-07-30

What's Changed

Added

- Made Flask-Limiter's `storage_uri` configurable via `RATELIMIT_STORAGE_URI`, defaulting to `"memory://"` (Flask-Limiter's own default — Render's behavior is unchanged). Fourth step of the Flask-to-Vercel migration (see `extras/deepdive_flask_to_vercel.md`): on Vercel, Fluid Compute's auto-scaling would otherwise give each concurrent instance its own in-memory rate-limit counter, silently multiplying every limit by however many instances are warm. Setting the env var to a `redis://` URL from an Upstash Redis database's "Connect" tab switches to shared, Redis-backed storage — confirmed by construction (`RedisStorage` selected when the env var is set, `MemoryStorage` when it's unset). Added `redis==7.4.1` to `requirements.txt` — pinned below 8.0.0, the ceiling limits (Flask-Limiter's dependency) requires for its `redis` extra. Getting the actual connection string remains a manual step: Vercel's own Marketplace docs for Upstash only surface a REST endpoint (`UPSTASH_REDIS_REST_URL/TOKEN`) for the JS `@upstash/redis` client, which Flask-Limiter can't use — the standard `redis:///redis://` string every Upstash database also exposes has to be copied from the Upstash console directly. (refs #38)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.21.1...v0.21.2>

v0.21.3 — 2026-07-30

What's Changed

Fixed

- Fixed `vercel.json`: an explicit `functions: {"app.py": {...}}` block disabled Vercel's zero-config Flask auto-detection entirely (build succeeded but produced no Serverless Function at all — every route 404'd). Replaced with `"framework": "flask"`, confirmed working end to end against a real preview deployment: `/`, `/_health`, `/form.js`, and `POST /submit` (multi-file upload, PDF build, ZIP bundling, Resend call, Redis-backed rate limiting) all verified against the actual Vercel runtime, not just docs. (refs #38)

Added

- Added `.vercelignore`. `vercel deploy` does not respect `.gitignore` for what gets uploaded — confirmed via `vercel deploy --dry --json` before ever running a real deploy, which showed `Clients/` (real client CVs, personal documents, 8.3MB), `extras/`, `.agents/`, `.claude/`, and the private report generators would otherwise all be uploaded to Vercel. Also excludes `.python-version` (pins 3.11.9 for local `pyenv`, conflicts with Vercel's Python 3.12 default) and `render.yaml/loading/` (Render-specific, unused on Vercel). (refs #38)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.21.2...v0.21.3>

v0.21.4 — 2026-07-30

What's Changed

Fixed

- `/submit` now always bundles the consultant's emailed attachment into a ZIP, even with zero client uploads — previously the ZIP was only built when the client attached at least one file, so the consultant received two different attachment formats depending on client behavior. Found while smoke-testing the Vercel migration (#38) with a full real-world submission (Alex Mercer's actual original intake data), but unrelated to the migration itself. `_build_zip()` already handled an empty uploads list correctly — the intake PDF alone still gets zipped. (closes #39)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.21.3...v0.21.4>

v0.21.5 — 2026-07-30

What's Changed

Added

- Added a "Testing the Intake Form" section to `SKILL.md`: use Alex Mercer's data (`Clients/Alex Mercer/`) for any form testing that needs realistic field values or file uploads — fictional, so no real-client privacy concern, and `Alex-Intake.pdf` already has a full set of original answers plus real upload files. Also documents the CSRF `Referer`-header requirement, the CV-or-fallback-fields validation rule, and the rate limit, discovered while smoke-testing the Vercel migration (#38).

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.21.4...v0.21.5>

v0.21.6 — 2026-07-30

What's Changed

Added

- Added a "Capturing development discussion in issue comments" policy to `SKILL.md`'s Versioning section, mirroring the `umoja-voices` project's own convention: post a `gh issue comment` summarizing real decisions, course-corrections, or noteworthy findings, in addition to (not instead of) the `CHANGELOG.md` entry, so the issue thread itself stays a readable record.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.21.5...v0.21.6>

v0.21.7 — 2026-07-30

What's Changed

Changed

- Untracked SKILL.md from git — it holds the consultant's plan-generation methodology and business know-how, kept local going forward rather than public on GitHub. Stays on disk unchanged; `git rm --cached` only, no content change.

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.21.6...v0.21.7>

v0.21.8 — 2026-07-30

What's Changed

Removed

- Removed `render.yaml` and `loading/index.html` — both Render services (`career-transition-intake`, `career-transition-loading`) have been deleted now that the Vercel migration (#38) is verified and cut over (landing page CTA now points at <https://career-transition-intake.vercel.app> in production). Stripped `/_health`'s CORS headers, which existed solely for the loading page's cross-origin polling — the endpoint itself stays, as a plain liveness probe. Updated `README.md` and the local `SKILL.md` to stop pointing at the dead loading-page URL and describe the current Vercel deployment instead of the old Render one. (closes #40)

Full Changelog: <https://github.com/ericgitonga/career-transition-intake/compare/v0.21.7...v0.21.8>

Part 2: career-transition — the landing page

Formation and evolution

The landing page was scaffolded directly as a Next.js (App Router, TypeScript, Tailwind CSS) app, matching the stack already used for `en-mascaradores` — a deliberate reuse of a known-good setup rather than a fresh technical decision. Its copy was adapted from a draft written against the "Build A Sales Page That Converts" framework (headline, pain points, offer, who-it's-for/who-it's-not, CTA, testimonials), originally drafted in the intake repo's own `Clients/Eric/eric_copy.md` and rewritten into the page itself. From the start its "Start My Plan" call-to-action pointed at the intake form's public entry point, hosted on Render — the two apps were always conceptually one product, just split across two repos and two hosting platforms from day one.

Its own independent history is much shorter than the intake app's: semantic versioning and a changelog were introduced immediately (mirroring the convention already established in the intake repo and in `umoja-voices`), followed by a Playwright E2E smoke suite gated on every PR, and — right at the point this document's scope ends — a pricing section. Everything after that (pointing the CTA at the newly-migrated Vercel intake app, fixing the `.vercelignore` monorepo issue, restyling the intake form to match, and repositioning the marketing copy toward comprehensiveness over speed) is post-merger work, tracked in the merged repo's own issue tracker rather than here.

Complete commit history

7 commits, oldest to newest:

- 62b6d199 (2026-07-28, ericgitonga) — Scaffold career-transition landing page
- 63a78ecd (2026-07-29, ericgitonga) — Introduce semantic versioning and CHANGELOG (v0.1.0, closes #2)
- af610e3a (2026-07-29, Eric Gitonga) — Merge pull request #3 from ericgitonga/versioning-changelog
- 36466030 (2026-07-29, ericgitonga) — Add Playwright E2E smoke suite gated on PRs to main (v0.2.0, closes #1)
- f07e2c74 (2026-07-29, Eric Gitonga) — Merge pull request #4 from ericgitonga/e2e-suite
- 2b263a86 (2026-07-30, Eric Gitonga) — Add pricing section with 50% introductory offer (#5)
- 678c990e (2026-07-30, Eric Gitonga) — Point intake CTA at the new Vercel-hosted Flask app (#6)

Issues (2 tickets, all comments in full)

Issue #1: Add Playwright E2E smoke suite gated on PRs to main

Opened by ericgitonga on 2026-07-28 — CLOSED — closed 2026-07-29

Set up an end-to-end regression suite mirroring the pattern used in umoja-voices/ebc-songs/merch-mockup (plain Python \playwright\ specs discovered/run by \e2e/run.py\, gated in CI on every PR to \main\ via \.github/workflows/e2e.yml\).

career-transition has no auth/DB/forms — it's a static marketing page — so this is the simplest of the four: cover the golden path (page loads, metadata correct, every section renders, footer copyright present) and confirm every "Start My Plan" CTA points at the intake form (<https://career-transition-loading.onrender.com>) and the contact link is a correct \mailto:\.

CI runs the suite against a production build (\next build && next start\) rather than a hand-run dev server — no Vercel/Supabase secrets needed since there's no database or serverless functions to diverge from a real deployment.

Issue #2: Introduce semantic versioning and CHANGELOG

Opened by ericgitonga on 2026-07-29 — CLOSED — closed 2026-07-29

Add \VERSION\, \CHANGELOG.md\, and a \version\ field in \package.json\, following the convention used in \umoja-voices\lebc-songs\lemerch-mockup\: Keep a Changelog format, SemVer (pre-1.0), a \tag: \vX.Y.Z\ line per entry, one real git tag + GitHub Release per version bump, and a 1:1 mapping between releases and the issue(s) they close.

Also add a local (gitignored) \ONBOARDING.md\ documenting the full workflow: issue-per-change, branch+PR (no direct pushes to \main\, matching the repo's existing branch-protection ruleset), CI as a merge gate, and the version-bump sequence — mirroring \merch-mockup\'s \ONBOARDING.md\.

Pull requests (4, with review comments)

PR #3: Introduce semantic versioning and CHANGELOG

Opened by ericgitonga on 2026-07-29 — MERGED — closed 2026-07-29

Closes #2.

Adds VERSION + CHANGELOG.md following the umoja-voices/ebc-songs/merch-mockup convention (Keep a Changelog + SemVer, pre-1.0, one tag+release per version bump). package.json's version field already read 0.1.0.

No CI to wait on for this PR (the e2e workflow lands in the next PR, #1) — pure docs/tooling addition.

PR #4: Add Playwright E2E smoke suite gated on PRs to main

Opened by ericgitonga on 2026-07-29 — MERGED — closed 2026-07-29

Closes #1.

Mirrors the umoja-voices/ebc-songs/merch-mockup e2e/ pattern (plain playwright specs discovered/run by e2e/run.py), simplified since this app has no auth/database/forms: golden-path page load + section rendering, and every CTA/contact link correctness check. CI (this PR's own .github/workflows/e2e.yml) runs the suite against a production build on the runner.

PR #5: Add pricing section with 50% introductory offer

Opened by ericgitonga on 2026-07-30 — MERGED — closed 2026-07-30

Summary

- Adds a "Simple, upfront pricing" section between How it works and Testimonials
- Shows KES 15,000 struck through against the introductory KES 7,500 (50% off)
- Notes the fee is paid when the intake form is submitted
- Adds an urgency line noting the introductory price won't last

Test plan

- [x] `npx tsc --noEmit` passes
- [x] Verified rendered section via `next dev` + SSR HTML inspection

□ Generated with [Claude Code](#)

PR #6: Point intake CTA at the new Vercel-hosted Flask app

Opened by ericgitonga on 2026-07-30 — MERGED — closed 2026-07-30

Summary

- Points the "Start My Plan" CTA at `career-transition-intake.vercel.app` instead of the Render loading page

- career-transition-intake has been migrated off Render onto Vercel (see career-transition-intake#38)
- Fluid Compute's cold starts are sub-second, so the loading page's cold-start workaround is no longer needed

Test plan

- [x] `npx tsc --noEmit` passes
- [x] Verified all 5 e2e tests pass locally against the updated URL
- [x] Verified the new production URL directly (`/` and `/_health` both 200)

□ Generated with [Claude Code](#)

Releases

2 releases:

v0.1.0 — 2026-07-29

Added

- Initial scaffold of the Career Transition Planning landing page: Next.js (App Router, TypeScript, Tailwind CSS), same stack as en-mascaradores. Hero, pain points, deliverables, who-it's-for/who-it's-not, how-it-works, testimonials, and contact sections, with the primary CTA linking out to the client intake form hosted separately at `career-transition-loading.onrender.com`.
- Semantic versioning and this changelog, following the convention used in `umoja-voices/ebc-songs/merch-mockup` (closes #2).

v0.2.0 — 2026-07-29

Added

- Playwright-based E2E smoke suite in `e2e/` covering the golden path (page loads, metadata correct, every section renders, footer copyright present) and every "Start My Plan" CTA / contact link pointing at the right target — mirroring the pattern used in `umoja-voices/ebc-songs/merch-mockup`, simplified since this app has no auth/database/forms of its own.
- `.github/workflows/e2e.yml` gates every PR to main on the suite, running it against a production build (`next build && next start`) on the runner — no Vercel/Supabase secrets needed since there's nothing here to diverge from a real deployment. (closes #1)